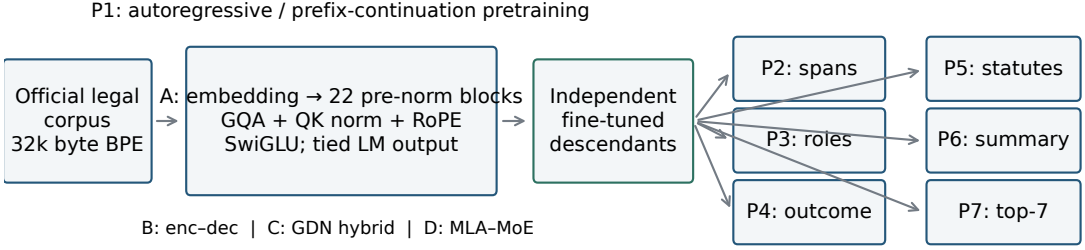
P1: autoregressive / prefix-continuation pretraining
Official legal corpus 32k byte BPE
A: embedding → 22 pre-norm blocks GQA + QK norm + RoPE SwiGLU; tied LM output
Independent fine-tuned descendants
P2: spans
P3: roles
P4: outcome
P5: statutes
P6: summary
P7: top-7
B: enc–dec | C: GDN hybrid | D: MLA–MoE